

STEP 1. UNDERSTAND DATASET

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns
```

```
In [2]: file = pd.read_csv("train_loan.csv")
file
```

Out[2]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	Coappli
0	LP001002	Male	No	0	Graduate	No	5849	
1	LP001003	Male	Yes	1	Graduate	No	4583	
2	LP001005	Male	Yes	0	Graduate	Yes	3000	
3	LP001006	Male	Yes	0	Not Graduate	No	2583	
4	LP001008	Male	No	0	Graduate	No	6000	
...
609	LP002978	Female	No	0	Graduate	No	2900	
610	LP002979	Male	Yes	3+	Graduate	No	4106	
611	LP002983	Male	Yes	1	Graduate	No	8072	
612	LP002984	Male	Yes	2	Graduate	No	7583	
613	LP002990	Female	No	0	Graduate	Yes	4583	

614 rows × 13 columns



```
In [3]: file.head()
```

Out[3]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	Coapplica
0	LP001002	Male	No	0	Graduate	No	5849	
1	LP001003	Male	Yes	1	Graduate	No	4583	
2	LP001005	Male	Yes	0	Graduate	Yes	3000	
3	LP001006	Male	Yes	0	Not Graduate	No	2583	
4	LP001008	Male	No	0	Graduate	No	6000	



```
In [4]: file.shape
```

Out[4]: (614, 13)

In [5]: `file.columns`

Out[5]: Index(['Loan_ID', 'Gender', 'Married', 'Dependents', 'Education', 'Self_Employed', 'ApplicantIncome', 'CoapplicantIncome', 'LoanAmount', 'Loan_Amount_Term', 'Credit_History', 'Property_Area', 'Loan_Status'], dtype='object')

In [6]: `type(file)`

Out[6]: `pandas.core.frame.DataFrame`

In [7]: `file.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 614 entries, 0 to 613
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Loan_ID               614 non-null   object
1   Gender                601 non-null   object
2   Married               611 non-null   object
3   Dependents            599 non-null   object
4   Education              614 non-null   object
5   Self_Employed         582 non-null   object
6   ApplicantIncome        614 non-null   int64
7   CoapplicantIncome      614 non-null   float64
8   LoanAmount             592 non-null   float64
9   Loan_Amount_Term       600 non-null   float64
10  Credit_History         564 non-null   float64
11  Property_Area          614 non-null   object
12  Loan_Status            614 non-null   object
dtypes: float64(4), int64(1), object(8)
memory usage: 62.5+ KB
```

In [8]: `file.value_counts()`

```

LP001005  Male    Yes    0      Graduate    Yes    3000
0.0      66.0      360.0      1.0      Urban
Y      1
LP002347  Male    Yes    0      Graduate    No     3246
1417.0    138.0    360.0      1.0      Semiurban
Y      1
LP002345  Male    Yes    0      Graduate    No     1025
2773.0    112.0    360.0      1.0      Rural
Y      1
LP002342  Male    Yes    2      Graduate    Yes    1600
20000.0   239.0    360.0      1.0      Urban
N      1

..
LP001674  Male    Yes    1      Not Graduate No     2600
2500.0    90.0     360.0      1.0      Semiurban
Y      1
LP001673  Male    No     0      Graduate    Yes    11000
0.0      83.0     360.0      1.0      Urban
..

```

In [9]: `file.describe()`

Out[9]:

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History
count	614.000000	614.000000	592.000000	600.000000	564.000000
mean	5403.459283	1621.245798	146.412162	342.000000	0.842199
std	6109.041673	2926.248369	85.587325	65.12041	0.364878
min	150.000000	0.000000	9.000000	12.000000	0.000000
25%	2877.500000	0.000000	100.000000	360.000000	1.000000
50%	3812.500000	1188.500000	128.000000	360.000000	1.000000
75%	5795.000000	2297.250000	168.000000	360.000000	1.000000
max	81000.000000	41667.000000	700.000000	480.000000	1.000000

STEP 2. DATA CLEANING

In [10]: `file.replace({"Loan_Status":{"N":0,'Y':1}},inplace=True)`

In [11]: `file.head()`

Out[11]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome
0	LP001002	Male	No	0	Graduate	No	5849	
1	LP001003	Male	Yes	1	Graduate	No	4583	
2	LP001005	Male	Yes	0	Graduate	Yes	3000	
3	LP001006	Male	Yes	0	Not Graduate	No	2583	
4	LP001008	Male	No	0	Graduate	No	6000	

In [12]: `file['Dependents'].value_counts()`

Out[12]:

```
0      345
1      102
2      101
3+       51
Name: Dependents, dtype: int64
```

In [13]: `loan = file.replace(to_replace='3+',value=3)`

In [14]: `loan.head()`

Out[14]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome
0	LP001002	Male	No	0	Graduate	No	5849	
1	LP001003	Male	Yes	1	Graduate	No	4583	
2	LP001005	Male	Yes	0	Graduate	Yes	3000	
3	LP001006	Male	Yes	0	Not Graduate	No	2583	
4	LP001008	Male	No	0	Graduate	No	6000	

In [15]: *convert categorical columns to numerical values*

```
loan.replace({'Married':{'No':0,'Yes':1}, 'Gender':{'Male':1,'Female':0}, 'Self_Employed':{'Not':0,'Yes':1}, 'Property_Area':{'Rural':0,'Semiurban':1,'Urban':2}, 'Education':{'Not Graduate':0,'Graduate':1}}
```

In [16]: `loan.head()`

Out[16]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome
0	LP001002	1.0	0.0	0	1	0.0	5849	
1	LP001003	1.0	1.0	1	1	0.0	4583	
2	LP001005	1.0	1.0	0	1	1.0	3000	
3	LP001006	1.0	1.0	0	0	0.0	2583	
4	LP001008	1.0	0.0	0	1	0.0	6000	

In [62]: `e = loan.drop(columns=['Loan_ID'],axis=1)`

In [63]: `e.head()`


Out[63]:

	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome
1	1.0	1.0	1	1	0.0	4583	1508.0
2	1.0	1.0	0	1	1.0	3000	0.0
3	1.0	1.0	0	0	0.0	2583	2358.0
4	1.0	0.0	0	1	0.0	6000	0.0
5	1.0	1.0	2	1	1.0	5417	4196.0

STEP 3. EXPLORATORY DATA ANALYSIS


```
In [64]: # Married & Loan Status
import seaborn as sns
sns.countplot(x='Married',hue='Loan_Status',data=loan)
```

Out[64]: <AxesSubplot:xlabel='Married', ylabel='count'>



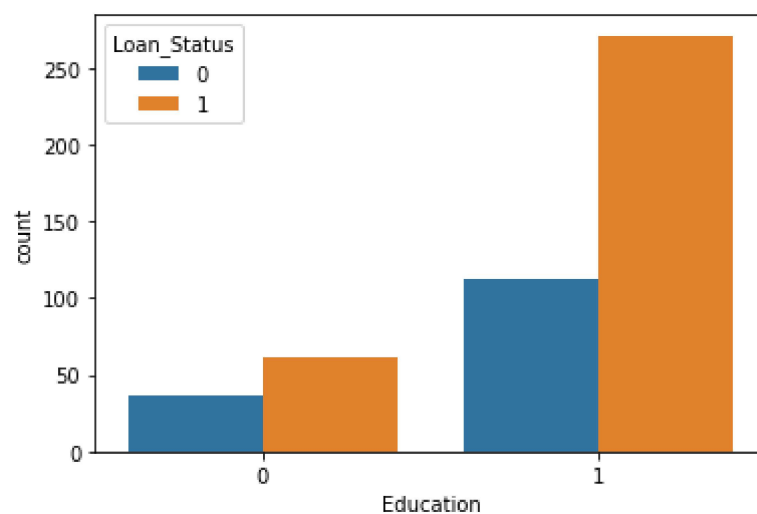
```
In [65]: sns.countplot(x='Dependents',hue='Loan_Status',data=loan)
```

Out[65]: <AxesSubplot:xlabel='Dependents', ylabel='count'>



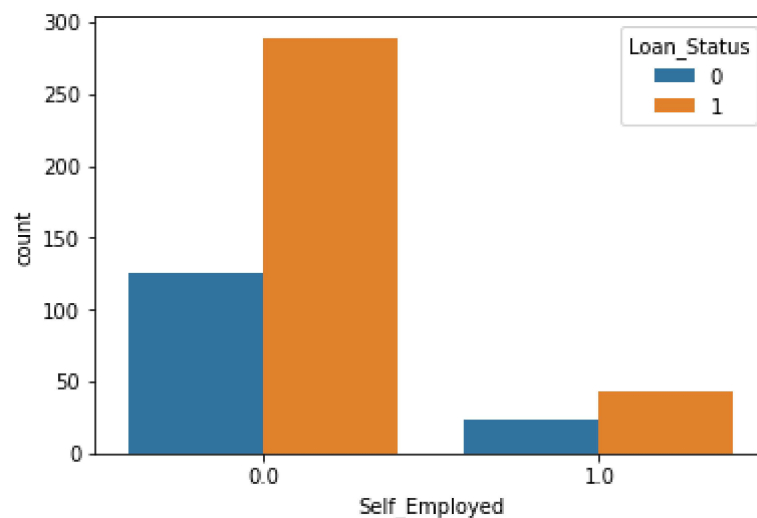
```
In [66]: sns.countplot(x='Education',hue='Loan_Status',data=loan)
```

```
Out[66]: <AxesSubplot:xlabel='Education', ylabel='count'>
```



```
In [67]: sns.countplot(x='Self_Employed',hue='Loan_Status',data=loan)
```

```
Out[67]: <AxesSubplot:xlabel='Self_Employed', ylabel='count'>
```



STEP 4. EXTRACT X AND y from the dataframe

```
In [78]: X = loan.drop(columns=['Loan_ID'],axis=1)
y = loan[['Loan_Status']]
```

```
In [79]: print(X)
print(y)
```

```
610      0      1
611      2      1
612      2      1
613      1      0
```

```
[480 rows x 12 columns]
```

```
Loan_Status
```

```
1      0
2      1
3      1
4      1
5      1
..     ...
609     1
610     1
611     1
612     1
613     0
```

```
[480 rows x 1 columns]
```

STEP 5. ONE HOT ENCODING

```
In [80]: encoding = pd.get_dummies(["Loan_Status"])
encoding
```

Out[80]:

Loan_Status	
0	1

STEP 6.MODEL BUILDING

```
In [86]: from sklearn.model_selection import train_test_split as tts
```

```
In [87]: X_train,X_test,Y_train,Y_test=tts(X,Y,test_size=0.3,stratify=Y,random_state=42)
```

```
In [93]: from sklearn.preprocessing import StandardScaler
scaler=StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```



```
In [96]: from sklearn.svm import LinearSVC
from sklearn.metrics import accuracy_score
model = LinearSVC()

# train the model on the training data
model.fit(X_train, Y_train)

# make predictions on the test data
y_pred = model.predict(X_test)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:63: Data ConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

return f(*args, **kwargs)

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\svm_base.py:985: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.

warnings.warn("Liblinear failed to converge, increase "

```
In [97]: accuracy = accuracy_score(Y_test, y_pred)
print('Accuracy:', accuracy)
```

Accuracy: 0.6944444444444444

```
In [98]: from sklearn.metrics import confusion_matrix

# calculate the confusion matrix
cm = confusion_matrix(Y_test, y_pred)

# print the confusion matrix
print(cm)
```

```
[[ 0 44]
 [ 0 100]]
```

```
In [99]: from sklearn.metrics import classification_report
# print the classification report
print(classification_report(Y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.00	0.00	0.00	44
1	0.69	1.00	0.82	100
accuracy			0.69	144
macro avg	0.35	0.50	0.41	144
weighted avg	0.48	0.69	0.57	144

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\metrics_classification.py:1248: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, msg_start, len(result))

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\metrics_classification.py:1248: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, msg_start, len(result))

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\metrics_classification.py:1248: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, msg_start, len(result))

STEP 7. PERFORMANCE COMPARISONS

```
In [101]: from sklearn.linear_model import LogisticRegression
# create a logistic regression model with default hyperparameters
lr_model = LogisticRegression()

# train the logistic regression model on the training data
lr_model.fit(X_train, Y_train)

# make predictions on the test data using the logistic regression model
lr_y_pred = lr_model.predict(X_test)

# calculate the accuracy of the logistic regression model
lr_accuracy = accuracy_score(Y_test, lr_y_pred)
print('Logistic regression accuracy:', lr_accuracy)

# print the classification report for the logistic regression model
print('Logistic regression classification report:')
print(classification_report(Y_test, lr_y_pred))
```

```
Logistic regression accuracy: 1.0
Logistic regression classification report:
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	44
1	1.00	1.00	1.00	100
accuracy			1.00	144
macro avg	1.00	1.00	1.00	144
weighted avg	1.00	1.00	1.00	144

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:63: Data ConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
return f(*args, **kwargs)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:763: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
In [102]: from sklearn.linear_model import SGDClassifier
# create a SGDClassifier model with default hyperparameters
sgd_model = SGDClassifier()

# train the SGDClassifier model on the training data
sgd_model.fit(X_train, Y_train)

# make predictions on the test data using the SGDClassifier model
sgd_y_pred = sgd_model.predict(X_test)

# calculate the accuracy of the SGDClassifier model
sgd_accuracy = accuracy_score(Y_test, sgd_y_pred)
print('SGDClassifier accuracy:', sgd_accuracy)

# print the classification report for the SGDClassifier model
print('SGDClassifier classification report:')
print(classification_report(Y_test, sgd_y_pred))
```

SGDClassifier accuracy: 0.4513888888888889

SGDClassifier classification report:

	precision	recall	f1-score	support
0	0.34	0.84	0.48	44
1	0.80	0.28	0.41	100
accuracy			0.45	144
macro avg	0.57	0.56	0.45	144
weighted avg	0.66	0.45	0.44	144

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:63: Data ConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
return f(*args, **kwargs)
```

```

In [103]: from sklearn.svm import LinearSVC, SVC
from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split

# Generate some synthetic data
X, y = make_classification(n_samples=1000, n_features=10, n_informative=5, random_s

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_s

# Train and evaluate a LinearSVC model
linear_svc = LinearSVC(random_state=42)
linear_svc.fit(X_train, y_train)
linear_svc_score = linear_svc.score(X_test, y_test)
print(f"LinearSVC score: {linear_svc_score:.3f}")

# Train and evaluate an SVC model with a linear kernel
svc_linear = SVC(kernel="linear", random_state=42)
svc_linear.fit(X_train, y_train)
svc_linear_score = svc_linear.score(X_test, y_test)
print(f"SVC with linear kernel score: {svc_linear_score:.3f}")

# Train and evaluate an SVC model with a polynomial kernel
svc_poly = SVC(kernel="poly", degree=3, random_state=42)
svc_poly.fit(X_train, y_train)
svc_poly_score = svc_poly.score(X_test, y_test)
print(f"SVC with polynomial kernel score: {svc_poly_score:.3f}")

# Train and evaluate an SVC model with an RBF kernel
svc_rbf = SVC(kernel="rbf", gamma=0.1, random_state=42)
svc_rbf.fit(X_train, y_train)
svc_rbf_score = svc_rbf.score(X_test, y_test)
print(f"SVC with RBF kernel score: {svc_rbf_score:.3f}")

# Train and evaluate an SVC model with a sigmoid kernel
svc_sigmoid = SVC(kernel="sigmoid", random_state=42)
svc_sigmoid.fit(X_train, y_train)
svc_sigmoid_score = svc_sigmoid.score(X_test, y_test)
print(f"SVC with sigmoid kernel score: {svc_sigmoid_score:.3f}")

```

```

LinearSVC score: 0.835
SVC with linear kernel score: 0.845
SVC with polynomial kernel score: 0.895
SVC with RBF kernel score: 0.935
SVC with sigmoid kernel score: 0.665

```

```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\svm\_base.py:985: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.
  warnings.warn("Liblinear failed to converge, increase "

```

In []:

